

UI PAGE SPECIFICATION

Voice AI Referral Conversion Platform

Page-by-page layout, components, states, and data — built to feed a one-shot UI prototype

Operational tool — not a marketing site • connector-first MVP scope • modeled on the supplied dashboard reference

Contents

1. How to use this spec

This document specifies every screen for the prototype UI. It is written to be handed directly to a one-shot UI build: each page section gives the layout structure, the components on it, the visual states it must show, and the data each component displays. Build the screens in the priority order in Section 3.

Scope discipline: this is the connector-first MVP. The MA confirms bookings in the practice's existing system; the platform's job is to surface AI-captured slots, call records, and analytics. The richer native approval queue is post-migration and is called out where relevant as out of MVP scope — design the placeholder, not the full flow.

The one rule for this UI

This is an operational tool, not a marketing dashboard. Every card on a daily-driver screen must answer either “what do I do next?” or “is the AI working?” If a card does neither, cut it. Reserve rich data visualization for the Analytics page, where exploration is the point.

2. Design system (modeled on the reference)

The supplied reference works because of disciplined density: many small purposeful cards, one accent color carrying the hierarchy, and generous whitespace between cards. Reproduce that discipline; do not reproduce decorative charts that drive no action.

Token	Spec	Notes
Layout shell	Fixed left sidebar nav + top bar + scrollable card-grid content	Same skeleton as the reference on every page
Accent color	One green/teal primary; a single warm accent reserved for urgency	Urgency color appears ONLY for escalations and alerts
Cards	Rounded corners, soft 1px border, white on a tinted page background	No drop shadows stacked on gradients; keep it flat and clean
Typography	One sans family, 2 weights; sentence case everywhere	No ALL CAPS, no Title Case in labels
State color coding	Consistent palette for referral/call states, reused on every screen	Define once in a legend; never recolor states per page
Density	Operational screens: compact. Analytics: room to breathe	Daily views prioritize scannability over beauty
Empty / loading	Every list and card needs an explicit empty and loading state	Prototype should show at least one empty state

Shared components (build once, reuse everywhere)

- Left sidebar navigation with the page list, active-item highlight, and a settings entry pinned to the bottom (mirror the reference).
- Top bar: global search, environment/practice selector, notification bell (drives the urgent-alerts count), user menu.
- State pill: a single component that renders any referral or call state in its consistent color. Used in every table and detail view.
- Patient peek: clicking a patient name anywhere opens the Referral Detail view — this deep-link is global.
- Transcript panel: reusable inline transcript with AI summary at top, audio scrubber, and turn-by-turn text. Appears on multiple pages.
- Metric card: label, big number, small delta vs. baseline, optional sparkline. The atomic unit of the dashboards.

3. Page inventory & build priority

Eleven screens. The priority column tells the one-shot build what to render fully vs. stub. If you are optimizing for a single high-impact prototype, build the five “Core” pages completely; the rest can be lighter.

#	Page	Primary user	Build priority
1	Landing / sign-in	All	Light (minimal)
2	Dashboard (home)	MA / Coordinator	Core — model on reference
3	Referral list	Coordinator / MA	Core
4	Referral detail view	All	Core
5	AI call log (Voice AI record)	MA / Coordinator	Core
6	Analytics dashboard	Administrator	Core
7	Urgent alerts / escalation queue	MA / Coordinator	High (trust/safety)
8	MA action queue (confirm slots)	MA	High (keystone flow)
9	Calendar	MA / Coordinator	Medium
10	Settings / configuration	Administrator	Medium (mostly static)
11	Referral ingestion / upload	Coordinator	Light (front door)

4. Page-by-page specification

Each page below gives: purpose, layout, the components on it, the states it must visually demonstrate, and the data each surface shows. Build in the order from Section 3.

4.1 Landing / sign-in

Purpose Internal entry point. This is a tool, not a marketing page — resist hero sections.

Layout Centered card on the tinted background: product mark, one-line descriptor, sign-in fields, sign-in button. Nothing else.

Components Logo lockup; email + password (or SSO button); primary sign-in button; subtle build/version footer.

States Default; invalid-credentials inline error; loading on submit.

Data None real — static. The prototype can route any sign-in straight to the Dashboard.

4.2 Dashboard (home) — model this on the reference

Purpose The MA's daily driver. Answers “what needs me right now” and “is the AI working” in one glance. This is the screen to make beautiful.

Layout Reference skeleton: sidebar + top bar + card grid. Top strip = urgent alerts (loud). Then a row of metric cards. Then two columns: action queue (left, wider) and pipeline + upcoming (right).

Cards on this page (in order)

- **1.** Urgent alerts strip — full-width banner at the very top. Count of escalations needing a human, each clickable into the escalation queue. If zero, a calm “all clear” state (still present, not hidden).
- **2.** Today's AI activity — row of metric cards: calls made, calls connected, slots captured, confirmations completed, no-answers. Each with delta vs. typical day.
- **3.** Needs your action — the AI-captured slots awaiting MA confirmation, as a compact list with patient, captured slot, and a confirm button. This is the bridge to page 8.
- **4.** Referral pipeline funnel — horizontal funnel: ingested → contacted → slot accepted → confirmed → completed, with counts. Operational, not decorative.
- **5.** Upcoming appointments peek — next few confirmed visits, read-only, with a link to the calendar.

States Populated; zero-urgent (calm strip); empty action queue (“nothing waiting on you”); loading skeletons per card.

Data Per-card counts for the day; the action queue items (patient, slot, transcript link); funnel stage counts; next appointments. All clickable through to detail.

Modeling note

The reference's strength is purposeful density with one accent. Keep the daily dashboard's charts minimal — the funnel is the only real visualization here. Push the satisfaction-score / rich-chart instinct to the Analytics page. Cards here earn their place by being actionable.

4.3 Referral list

Purpose The workhorse table. Where a coordinator/MA spends most of the day triaging and tracking referrals.

Layout Full-width data table under a filter/sort bar. Sticky header. Row click → Referral detail.

Columns

- Patient name · demographics (age/sex) · preferred language · insurance · referral reason · referring provider · referral time · location · current state (state pill) · priority · call attempts (count) · transcript (icon link if exists) · transcript summary (truncated, hover/expand).

Filters / sort Filter by state, priority, language, referral source, location; sort by referral time, attempts, priority. A prominent “urgent only” toggle.

States Populated; filtered-empty; loading; a high-priority row visually distinct; a row with no valid phone flagged.

Data One row per referral with every column above; summary text generated from the latest call transcript.

4.4 Referral detail view

Purpose The single source of truth for one patient. Everything links here. The most-opened screen after the list.

Layout Two-column: left = patient & referral facts; right = activity timeline + transcripts. Action bar pinned top-right.

Sections

- Patient & insurance panel — name, demographics, contact, preferred language, insurance/network, location.
- Referral facts — reason, referring provider, referral time, urgency/priority, current state pill.
- State timeline — vertical history: ingested → each outreach attempt → slot accepted → confirmed, timestamped, with outcomes.
- Call attempts — each attempt: time, channel, outcome, and an expandable transcript panel with AI summary + audio.
- Next scheduled action — what the system will do next and when; manual override controls (pause, escalate, change cadence).

States Active in cadence; slot accepted (awaiting MA); booked; escalated (banner); closed-won; closed-lost with reason + re-activate button.

Data Full patient + referral record; ordered attempt history; transcripts + summaries; next-action schedule; audit attribution.

4.5 AI call log (Voice AI record)

Purpose The audit and review surface for everything the voice AI did. Proves the AI is behaving and lets staff inspect any call.

Layout Filterable log table; selecting a row opens a side panel (or expands inline) with the transcript panel.

Columns / panel

- Row: timestamp · patient (click → referral detail) · attempt # · channel (voice/SMS) · outcome state pill · duration · outcome summary · transcript link.
- Panel: AI summary at top, audio scrubber, full turn-by-turn transcript, disclosure-played indicator, escalation flag if raised.

Filters By outcome (connected / voicemail / no answer / declined / escalated), by date, by patient, by language.

States Mixed-outcome populated log; escalated call visually flagged; voicemail row; no-transcript row; loading; empty (no calls yet today).

Data One row per call/SMS attempt with outcome + summary; transcript text; audio reference; disclosure + escalation metadata.

4.6 Analytics dashboard (administrator ROI)

Purpose The strategic, buyer-facing screen. Distinct from the daily dashboard — this is where rich visualization belongs. Makes the ROI undeniable.

Layout Reference-style card grid, but here charts are the point: a hero conversion funnel, KPI cards, breakdown charts, and the migration-readiness gauge.

Panels

- Conversion funnel with lift — referral → contacted → slot accepted → confirmed → completed, with conversion lift vs. pre-platform baseline.
- Recovered revenue — estimated dollars from converted + backlog-reactivated referrals; the headline KPI card.
- Backlog re-activation results — outcome of running the dead-referral pile; the first-week proof.
- Leakage by stage & source — where patients drop, broken down by referral source; the actionable diagnostic.
- Shadow-calendar fidelity — a gauge/trend of mirror correctness vs. the real schedule: the migration-readiness signal. Visually distinct (it is internal, not ROI).
- Compliance & audit snapshot — disclosures, quiet-hour adherence, opt-outs, change attribution; a summarized, exportable block.

States Populated with a healthy lift story; a low-fidelity warning state on the gauge; date-range selector; loading.

Data Funnel counts + baseline; revenue model output; backlog stats; leakage breakdown; fidelity time series; compliance counts.

Why fidelity gets its own panel

The shadow-calendar fidelity metric is not an ROI number — it is the instrument that governs the entire migration thesis. It must be visible and visually separated from the revenue story so an administrator (and your own team) can read migration-readiness at a glance.

4.7 Urgent alerts / escalation queue

Purpose The safety surface. Calls that went wrong, patient distress, clinical questions, bad data — anything needing a human now. Design it to feel urgent, never buried.

Layout High-contrast list using the reserved warm accent. Each item is a card with reason, context, and a single clear action.

Each escalation card

- Escalation reason (distress / clinical question / wrong number / bad data / call failure), patient (click → detail), timestamp, the triggering transcript inline, and an “I’m handling this” / reassign action.

States Active escalations (loud); claimed/in-progress (dimmed, shows owner); resolved (collapses out); empty = reassuring all-clear, still a real screen.

Data Escalation reason + severity; linked patient/referral; transcript excerpt; owner + status.

4.8 MA action queue (confirm captured slots)

Purpose The connector-mode keystone. AI captured a slot from a call → it lands here → MA confirms it in the practice system and marks done.

Layout Focused queue: one card per pending slot, transcript inline, decisive action buttons. Optimized for fast repetitive triage.

Each queue card

- Patient + referral summary, the AI-captured slot (date/time/provider), the call transcript + AI summary inline, and actions: confirm (mark booked in practice system), edit slot, reject (re-enter outreach).

States Pending items; one expanded with transcript; just-confirmed (animates out); empty (“nothing waiting”); loading.

Data Captured slot details; linked patient/referral; transcript + summary; confirmation result writes back to state + mirrors to shadow calendar (shown as a subtle status, no MA action).

MVP scope marker

This is the Connector-mode confirm queue — the MA commits in their own system, the platform records it. The richer Native approval queue (AI writes onto our calendar, MA approves there) is post-migration and out of MVP scope. Design a labeled “v2” placeholder if you want it visible, but do not build the native approval flow.

4.9 Calendar

Purpose See scheduled appointments. In the MVP this is largely a read view reflecting confirmed bookings (and, for admin visibility, the shadow-mirror state).

Layout Day/week toggle, standard calendar grid, color-coded by state using the shared state palette.

Components Calendar grid; state legend; appointment chip (click → referral detail); day/week switcher; provider/location filter.

States Populated week; a confirmed vs. shadow-mirrored entry visually distinguishable; empty day; loading.

Data Confirmed appointments with patient, provider, location, state; read-only in MVP.

4.10 Settings / configuration

Purpose Where the “supervised automation, not a black box” story is made tangible: cadence policy, AI scripts & disclosure, quiet hours, escalation rules. Mostly static in the prototype but high narrative value.

Layout Left sub-nav of setting groups; right = the selected group's form. Reference's Settings nav item is the entry.

Setting groups

- Cadence policy — attempt schedule, channel escalation, stop rules.
- AI scripts & disclosure — editable script blocks, mandatory AI-disclosure language (read-only/locked look), language variants.
- Quiet hours & compliance — contact windows, opt-out handling, PHI-on-voicemail rule.
- Escalation rules — what triggers a human; thresholds.
- Practice profile — specialties, providers, locations, languages, hours.

States View vs. edit; saved confirmation; a locked/compliance-controlled field shown as non-editable.

Data Config values; can be static defaults in the prototype.

4.11 Referral ingestion / upload

Purpose The front door — getting referrals into the system. Manual upload first in the MVP.

Layout Upload dropzone + a preview/validation table of parsed rows before commit.

Components Drag-drop upload; parsed-rows preview with validation flags (bad phone, missing field, possible duplicate); confirm-import button; per-row fixable fields.

States Empty dropzone; parsing/loading; preview with some rows flagged; import success summary.

Data Parsed referral rows + validation status; on import they enter the referral list as Ingested.

5. Cross-cutting requirements

These apply to every screen and matter for a convincing prototype:

- **Consistency:** State color system — define the referral/call state palette once; apply identically on list, detail, call log, calendar, and dashboard. Inconsistent state colors break trust instantly.
- **Navigation:** Global patient deep-link — a patient name is clickable everywhere and always lands on the Referral detail view. Wire this on every table and card.
- **States:** Every list/card needs an explicit empty and loading state — the prototype should show at least one empty state on purpose; it makes the product feel real.
- **Reuse:** Transcript panel is one reusable component — same look on detail, call log, escalation, and action queue. Build it once.
- **Signal:** Urgency is a reserved color — the warm accent appears only for escalations/alerts, nowhere decorative. It must mean “a human is needed.”
- **Trust:** Audit visibility — confirmed/changed-by attribution is shown on detail and call log even in the MVP; it is the spine of the eventual native migration.

6. Suggested one-shot build prompt structure

When you feed this to a one-shot UI generator, structure the request as: (1) the design-system block from Section 2, (2) the shared-components list, (3) the five Core pages fully specified from Section 4, (4) the cross-cutting rules from Section 5, then (5) the remaining pages as lighter specs. Lead with the Dashboard and Analytics screens since they carry the demo, and explicitly instruct: operational tool, reference-modeled, one accent color, no decorative cards on daily views.

— End of specification —